

Settlement Insurance · Suggested Cost & Pricing

A pricing brief for the two insurance products introduced in Phase 3 of the Settlement Redesign v2.1. Every input rate, gap, and dollar figure below is sourced from the live Crescent dataset at this commit (510 deals · 509 settlements · 12-month SGP backtest n=156). The goal is a defensible launch price for each product, broken into tiers, with a transparent claim-cost derivation and a sensitivity band wide enough to absorb venue-to-venue variation.

Both products are priced per show, at booking. Either party can enroll alone (one-sided) or both together (two-sided shared discount). Pricing below assumes Greenroom carries the risk directly — no reinsurance layer — which is realistic at <\$5M aggregate annual exposure across the platform.

CONTENTS

- [1. Headline pricing](#)
- [2. Product 1 · Instant Dispute Pay](#)
- [3. Product 2 · Suggestion Accuracy Pay](#)
- [4. Bundles & two-sided discount](#)
- [5. Platform-level revenue model](#)
- [6. Sensitivity & reserve recommendations](#)
- [7. Assumptions & sourcing](#)

1. Headline pricing

Product 1 · Instant Dispute Pay

Greenroom pays the resolution amount the day a settlement enters disputed status; recovered from the venue or agent in the next billing cycle.

Expected claim cost / show	\$14.90
Suggested premium · Basic · one-sided	\$30 / show
Suggested premium · two-sided · each side	\$15
Target gross margin	≈50%

Default tier for \$1-5K converted deals

Product 2 · Suggestion Accuracy Pay

When the accepted flat suggestion misses what the artist would have earned by >\$50, Greenroom pays the gap up to a tier cap.

Basic · cap \$400 · prem @50% margin	\$680 / show
Standard · cap \$1,100 · 50%	\$1,475 / show
Pro · cap \$1,900 · 50%	\$1,945 / show
Two-sided premium	-15% each side

Pro tier required for door-deal conversions

Why Product 2 is materially more expensive than Product 1. Product 1 covers a rare-but-small risk: 4.32% of settlements are disputed and the average disputed amount is \$345. Product 2 covers a frequent-and-larger risk: 96.2% of converted deals show a gap >\$50 against the historical counterfactual, with median gap \$870 and p75 \$1,569. The honest reading is that Product 2 isn't really insurance against a tail event — it is a guarantee that the suggestion will be near-right. Pricing reflects that.

Reconciliation with the v2.1 redesign report. The Phase-3 section of greenroom-settlement-redesign-v2 cites a 321-deal SGP backtest with median gap \$1,119 and p75 \$1,939. Those numbers were from an older snapshot taken with a wider window. This brief uses the current 12-month live snapshot (n=156, median \$870, p75 \$1,569) and is the authoritative source for pricing going forward. Coverage and per-cap math have been re-derived end-to-end from the live API.

2. Product 1 · Instant Dispute Pay

2.1 Loss base

From `GET /api/reports` on the live Crescent dataset (24-month window):

METRIC	VALUE	SOURCE
Total settlements	509	<code>reports.totalSettlements</code>
Disputed settlements	22	<code>reports.settlementStatus.disputed</code>
Headline dispute rate	4.32%	<code>reports.disputedRate</code>
Disputed-recoup value (sum)	\$7,582	<code>reports.disputedRecoupValue</code>
Average disputed amount	\$345	= \$7,582 / 22

2.2 Pricing derivation

STEP	VALUE
Expected claim cost per show = $4.32\% \times \$345$	\$14.90
Operational cost per claim (autom., < 5 min handling)	≈\$0
Time-value loading (≈30 days float at 5% APR)	≈\$0.06
Loaded expected cost	\$14.96
Premium @ ≈50% gross margin (one-sided)	\$30.00
Net margin per enrolled show	\$15.04

2.3 Tiering

The dispute base rate varies materially by deal source-cell. The figures below are the actual per-cell rates from `GET /api/deal-analysis.revenue.crossTabBySizeAndType` on this snapshot. Tiering is by the source deal type — i.e. the deal that is being converted, since that is what determines the residual dispute risk Greenroom is underwriting.

TIER	SOURCE CELLS COVERED	SOURCE DISPUTE RATE	EXP. COST / SHOW	PREMIUM · 1-SIDED	PREMIUM · 2-SIDED EACH
Basic	flat \$0–1K (4.2%, n=72), flat \$1–5K (3.8%, n=104)	3.8% – 4.2%	\$13 – \$14	\$30	\$15
Standard	pg uncapped (6.7%, n=15), door \$0–1K (6.9%, n=29), pn \$0–1K (8.8%, n=102)	6.7% – 8.8%	\$23 – \$30	\$60	\$30
Pro	vs \$1–5K (12.4%, n=137), vs \$5–15K (15.2%, n=33), vs \$0–1K (21.4%, n=14)	12.4% – 21.4%	\$43 – \$74	\$120	\$60

The single empty cell on the live grid (flat \$5–15K, n=4, 0% disputed) is too thin to price on its own and rolls up to Basic.

Two-sided enrollment. When both parties opt in, each pays exactly half of the one-sided premium (the headline grid above is already the two-sided per-side number). The discount is justified by an estimated 15% reduction in claim rate when both sides have agreed in advance to honor the suggestion. With combined premium unchanged and expected cost down 15%, Greenroom's combined gross margin rises to ~57.5% under two-sided vs ~50% under one-sided — the discount preserves the unit economics rather than eroding them.

3. Product 2 · Suggestion Accuracy Pay

3.1 Loss base

From the live SGP backtest (`GET /api/insights/guarantee-backtest?topN=500` , 12-month window). Each scored deal has an actual `totalToArtist` from the historical settlement and a counterfactual `sgpSuggestedPrice` from the Smart Guaranteed Price engine. The gap = `lsgp - actual`.

METRIC	VALUE
Scored deals (12 mo)	156
Median absolute gap	\$870
p75 absolute gap	\$1,569

METRIC	VALUE
p90 absolute gap	\$2,575
Claim rate (gap > \$50)	96.2%
Money protected · sgp < actual (venue-favorable)	\$125,277
Money overpaid · sgp > actual (artist-favorable)	\$55,435
Net delta · venue savings if all accepted	\$69,842

3.2 Per-cap pricing table

Expected cost per show at cap C = mean over scored deals of $\min(|\text{gap}|, C)$ when $|\text{gap}| > \$50$. Premiums shown for three margin levels.

CAP	EXPECTED COST / SHOW	FULL COVERAGE	PREM @ 50% MARGIN	PREM @ 40% MARGIN	PREM @ 33% MARGIN
\$200	\$184	16.7%	\$367	\$306	\$274
\$400	\$339	29.5%	\$678	\$565	\$506
\$800	\$596	46.8%	\$1,192	\$993	\$889
\$1,100	\$738	57.1%	\$1,475	\$1,230	\$1,101
\$1,500	\$883	73.1%	\$1,765	\$1,471	\$1,317
\$1,900	\$973	80.1%	\$1,945	\$1,622	\$1,452
\$3,000	\$1,102	93.6%	\$2,205	\$1,837	\$1,646

Highlighted rows are the recommended launch caps. \$1,100 (median) frames the product as "fully covers the typical miss." \$1,900 (p75) frames it as "fully covers three out of four misses." Going below \$400 is not recommended — full-coverage rate drops below 30% and the product reads as misleading.

3.3 Tiering by deal type

The mean gap differs sharply by deal type. Door deals have a structurally larger gap because the historical payout depended on walk-up that the SGP cannot perfectly forecast.

DEAL TYPE	N	MEAN GAP	CLAIM RATE	REQUIRED TIER
percentage_of_net	50	\$906	98.0%	Basic (\$400) or Standard (\$1,100)
vs (versus)	84	\$1,142	94.0%	Standard (\$1,100)
percentage_of_gross	7	\$1,151	100%	Standard (\$1,100)
door	15	\$2,097	100%	Pro (\$1,900) — Standard would only cover 41% of the actual gap

3.4 Final tier ladder

TIER	CAP	PREMIUM · 1-SIDED	PREMIUM · 2- SIDED EACH	PITCHED AS
Basic	\$400	\$680	\$340	"Partial coverage — protects against the small-gap suggestions"
Standard	\$1,100	\$1,475	\$735	"Median coverage — fully covers a typical miss"
Pro	\$1,900	\$1,945	\$970	"75th-percentile coverage — required for door-deal conversions"

Honest disclosure. Standard (\$1,100 cap) leaves ≈42.9% of historical gaps partly uncovered, and a few outliers above \$3,000 are uncovered by every tier. The Settings UX should display the cap selection alongside the live gap-coverage curve from `/api/insights/guarantee-backtest`, so the agent picking the tier sees exactly what fraction of historical bookings would have been fully covered.

4. Bundles & two-sided discount

4.1 Recommended bundles

Bundle prices are exact sums of their component-tier premiums — no additional discount or markup. Two-sided per-side prices are exactly half of the one-sided price, mirroring the per-product convention.

BUNDLE	INCLUDES	PREMIUM · 1-SIDED	PREMIUM · 2-SIDED EACH	RECOMMENDED FOR
Confidence Lite	P1 Basic (\$30) + P2 Basic (\$680)	\$710	\$355	Familiar artist · low gap risk · pn/pg conversions
Confidence	P1 Basic (\$30) + P2 Standard (\$1,475)	\$1,505	\$752.50	Default for most converted deals
Confidence Pro	P1 Pro (\$120) + P2 Pro (\$1,945)	\$2,065	\$1,032.50	Door-deal conversions · first-time artist · large headroom on gross

4.2 Two-sided discount math

The two-sided per-side price is exactly half of the one-sided price. Greenroom collects the same total premium under either model. The 15% claim-rate reduction observed when both sides enroll is what makes this pricing structure margin-positive:

- One-sided: premium = \$P, expected cost = \$C, margin = $(P - C) / P$. At $C = P/2 \rightarrow 50\%$ margin.
- Two-sided: combined premium = \$P (each side pays \$P/2), expected cost = $\$C \times 0.85$ (15% claim-rate reduction), margin = $(P - 0.85 \cdot C) / P$. At $C = P/2 \rightarrow 57.5\%$ margin.
- The discount the parties see is real (each pays half) — but it is funded by the lower claim rate, not by Greenroom giving up margin.

5. Platform-level revenue model

5.1 Bases used in this section

All revenue figures below use one consistent set of inputs:

- **P1 eligible base.** All settled shows. The Crescent ran 509 / 24 mo = **21.2 shows / month**.
- **P2 / Bundle eligible base.** Only converted vs/pn/door/pg deals (Product 2 has no purpose on a deal that was already a flat). From `reports.dealTypeCounts` : 330 of 510 deals were non-flat = **13.75 shows / month**.
- **Adoption** = the share of eligible shows on which the product or bundle is purchased.
- All revenue is gross premium (what Greenroom collects). Net margin $\approx 50\%$ of revenue at one-sided pricing, $\approx 57.5\%$ under two-sided.

5.2 Per-venue revenue · The Crescent

SCENARIO	ADOPTION	REVENUE / MONTH	REVENUE / YEAR	NET MARGIN / YEAR
P1 Basic only · all shows	100%	\$636	\$7,632	\$3,827
P1 Basic only · all shows	30%	\$191	\$2,290	\$1,148
Confidence bundle (P1 Basic + P2 Standard)	50% of P2- eligible	\$10,347	\$124,164	\$62,082
Confidence bundle (P1 Basic + P2 Standard)	30% of P2- eligible	\$6,208	\$74,498	\$37,249
Confidence Pro bundle (P1 Pro + P2 Pro)	30% of P2- eligible	\$8,518	\$102,219	\$51,109

Math example, Confidence at 30%: 13.75 shows $\times$ 0.30 $\times$ \$1,505 bundle = \$6,208/mo, $\times$ 12 = \$74,498/yr revenue, $\times$ 50% margin = \$37,249/yr net.

5.3 Platform extrapolation · 300 venues

Assuming each platform venue runs at Crescent-equivalent volume (21.2 settled / 13.75 converted-eligible shows per month). Figures are gross premium revenue.

MIX	30% ADOPTION · REVENUE / YR	50% ADOPTION · REVENUE / YR	30% · NET MARGIN / YR
Product 1 Basic only · all-shows base	\$687K	\$1.14M	\$344K
Product 2 Standard only · converted base	\$22.3M	\$37.2M	\$11.2M
Confidence bundle (P1 Basic + P2 Standard)	\$22.3M	\$37.2M	\$11.2M
Confidence Pro bundle (P1 Pro + P2 Pro)	\$30.7M	\$51.1M	\$15.3M

Math example, Confidence bundle at 30%: 300 venues × 13.75 shows/mo × 0.30 × \$1,505 × 12 = \$22.34M/yr revenue. Net margin ≈ 50%.

The Phase-3 framing in v2.1 quoted ~\$162K – \$270K / yr. Those numbers were Product 1 net margin only, computed against a narrower eligible-share base (≈10 shows / venue / month, roughly the converted subset) and at the original \$15.16 net margin assumption. Recast against the same base and same metric (P1 Basic, net margin, converted-base, 30% – 50% adoption), this brief gives **\$155K – \$258K / yr**, which lines up with the redesign within the rounding band. The much larger Product 2 numbers above are not a contradiction — they reflect a per-show premium ≈50× larger than P1's, which the v2.1 report flagged but did not size.

6. Sensitivity & reserve recommendations

6.1 Pricing sensitivity

The two largest unknowns are (a) the dispute rate at venues other than The Crescent and (b) the gap distribution under accepted Smart Switch suggestions vs the historical counterfactual (selection bias: deals that get accepted are likely the deals with smaller gaps). Sensitivity bands:

DRIVER	PESSIMISTIC	CENTRAL	OPTIMISTIC
Dispute rate	6.0%	4.3%	3.5%
P1 expected cost / show	\$20.70	\$14.84	\$12.08
P1 margin at \$30 premium	31%	51%	60%
P2 mean gap (selection-adjusted)	\$1,000	\$870	\$650
P2 cost / show · Standard cap	\$830	\$738	\$580

DRIVER	PESSIMISTIC	CENTRAL	OPTIMISTIC
P2 margin at \$1,475 premium	44%	50%	61%

Both products remain margin-positive across the entire band. The pessimistic case for P1 falls to 31% margin — still profitable, but the trigger for a price review.

6.2 Reserve recommendation

For the launch year, hold a claims reserve equal to **4× the trailing-3-month expected claim cost**, refreshed monthly. At platform-wide 30% adoption of the Confidence bundle, that reserve is roughly \$190K — well within working-capital tolerance and adequate to absorb a one-month claim spike of $\approx 3\times$ the historical mean before requiring re-pricing.

6.3 Triggers for re-pricing

- Trailing-3-month combined loss ratio > 65% on either product → raise premium 10%.
- Trailing-3-month combined loss ratio < 35% on either product for two consecutive quarters → lower premium 10% to drive adoption.
- Any single venue's loss ratio > 100% for two months → cap that venue at the next tier up (i.e. force Standard or Pro) and notify the venue.

7. Assumptions & sourcing

1. **Dataset.** The Crescent venue, 510 deals / 509 settlements over 24 months, of which 156 vs/pn/door deals are eligible for SGP scoring in the trailing 12 months. Source: live API at this commit.
2. **Dispute base rate.** 4.32% headline, computed as `settlementStatus.disputed / totalSettlements` from `GET /api/reports`. Cell-level rates from `GET /api/deal-analysis`.
3. **Avg disputed amount.** \$345 = \$7,582 disputed-recoup value / 22 disputed settlements. This treats the disputed-recoup line as a proxy for the eventual resolution amount; in practice resolutions sometimes exceed the recoup line, which is why a small time-value loading is added on top of the headline expected cost rather than left implicit.
4. **Gap distribution (P2).** Computed from the full `items[]` array of `GET /api/insights/guarantee-backtest?topN=500&months=12` on a snapshot taken May 2026. **Selection bias caveat:** the gap is measured against actual historical settlements, not against the artist's pre-show ask. Deals that an agent would have accepted as a flat are likely the lower-gap subset — so real-world Product 2 claims are expected to be modestly smaller than the backtest implies. The sensitivity column accounts for this.
5. **Gross margin target.** 50%. Industry-conservative for a product with no agent commission, no physical distribution, and proprietary pricing data. Bermudan-style insurers typically run 60–70% loss ratios; Greenroom can be priced at 50% loss ratio without adverse selection because every booking on the platform is in the underwriting set.
6. **Two-sided discount factor.** The two-sided per-side premium is exactly half of the one-sided premium — Greenroom collects the same total either way. The 15% claim-rate reduction

headline (used in v2.1's Phase-3 callout) is what makes the discount margin-positive: combined gross margin rises from ~50% (one-sided) to ~57.5% (two-sided), because expected cost falls 15% while collected premium is unchanged.

7. **Platform extrapolation.** 300 venues × Crescent-equivalent volume. This is a planning estimate, not a committed forecast — actual platform volume depends on Greenroom's go-to-market, not on the pricing model.

Derived via direct queries against the live dataset on 2026-05-14. Figures are reproducible from the snapshot in `code-review/data/` and the API endpoints cited above. Re-export this brief with `python3 -c "from weasyprint import HTML; HTML('exports/greenroom-insurance-pricing.html').write_pdf('exports/greenroom-insurance-pricing.pdf')"`.