

Simpler Deals, Smarter Tools

Greenroom's settlement tool works for flat guarantees but breaks on three deal types. The fix is not to build more sophisticated analytics on top of a complex system — it is to simplify the underlying deal structures, then let the tools do more with less.

Every figure in this revision sourced from `/api/deal-analysis`, `/api/reports`, `/api/insights/switch-savings` & `/api/insights/switch-projected-grid` at this commit. 24-month window unless otherwise noted.

THE CRESCENT • LAST 24 MONTHS

SETTLED SHOWS 509 \$3.22M gross · \$1.94M to artists	SETTLEMENT DISPUTE RATE 4.3% 22 of 509 settlements	IN-APP TOOL COVERAGE 38% flat + % of gross only	SETTLE-NIGHT FRICTION 62% door / vs / % of net deals
--	---	---	---

*The right approach is to **simplify first, then automate**. Clean deal structures produce consistent, predictable data. Consistent data is what AI tools can actually work with — to suggest the right deal at booking, answer questions before show night, and reduce the manual entry that causes disputes. Trying to solve structural complexity with*

data cleaning gets harder every year. Simplifying the structure solves it once.

IMPLEMENTATION ROADMAP

PHASE 1

Fix door deals

93% of door deals leave the venue underwater. The structure pays the artist the entire pool above expenses by construction — the venue has no upside.

Immediate • MVP

PHASE 2

Simplify vs / % of net

In the \$1–5K bucket vs deals dispute at 12.4% vs 3.8% for equivalent flat. The percentage clause adds friction without changing the payout.

Short-term • MVP

PHASE 3

Insurance services

Offer the safety net that makes pre-agreed flats safe to commit to weeks before show night.

Post-MVP

PHASE 1 — DOOR DEALS

Phase 1 • MVP

Replace the door deal structure with a hybrid

29 shows • 24 months • all \$0–1K bucket

PROBLEM

Under the current door arrangement, the venue covers all production costs and the artist takes the entire pool above those expenses. By construction, the venue's net on a door deal is zero on a perfect-cost-recovery night and negative anywhere expenses run over. **Aggregate venue net across all 29 door deals in the last 24 months is $-\$29,056$** , an average of $-\$1,007$ per show.

The cross-tab is unambiguous: **27 of 29 door deals (93.1%) end with the venue's net negative**. No other deal-type cell at this venue has a loss-making rate above 53%. There is no upside case in the historical data — door

deals are a structurally one-sided arrangement at the gross levels The Crescent actually books.

FINDING

The clearest single example is Pale Lake ([show_0207](#) , 2025-03-22). Gross box office was \$19,296 with \$1,909 in expenses, the artist was paid the entire pool — \$17,387 — and the venue's net came in at exactly \$0. Under a hybrid structure with a \$500 artist floor, a 60/40 split, and a \$1,500 expense cap, the same night reorders the economics for both parties.

CURRENT — DOOR DEAL	PROPOSED — SMART SWITCH HYBRID
Pale Lake artist payout \$17,387	Pale Lake artist payout \$10,020
Pale Lake venue net \$0	Pale Lake venue net +\$7,367
Venue stake in upside None by construction	Venue stake in upside 40% above the cap
24-month venue net (29 shows) - \$29,056	24-month projected savings + \$50,220
Venue underwater on 27 of 29 (93.1%)	Loss-nights avoided 17 (27 → 10)

The Smart Switch counterfactual replays each historical door deal with the hybrid: \$500 floor, 60% of the pool above a \$1,500 expense cap. **Across all 29 shows, projected venue savings total \$50,220 — by far the single largest contribution to the "If Smart Switch had been used" rollup on the Insights tab**, and more than double the next-largest cell (vs / \$1–5K).

This is not a model artefact. The hybrid keeps the artist's economics meaningful (a \$500 floor on slow nights, plus 60% of any real walk-up); it just lets the venue stop eating cost overruns on the 27 of 29 nights where today it absorbs the entire downside.

SUGGESTION

At deal entry, Greenroom proposes hybrid terms instead of a pure door arrangement: **\$500 artist floor + 60% / 40% split of the pool above a \$1,500 expense cap**. Both parties see the projected outcome at the cell-mean gross of \$4,420 before committing. For expected gross above \$15,000 the engine

intentionally suppresses a suggestion (insufficient historical data at that level for The Crescent) and the parties negotiate directly.

The AI layer (Phase 2 and beyond) will be able to explain this structure in plain language to agents and answer "what would the artist earn if the show does \$X?" before the deal is signed — exactly the kind of question that currently gets resolved at midnight.

EXPECTED OUTCOME

The venue participates in the upside of a strong show for the first time. The artist retains a meaningful share of every viable night. Both door-deal disputes in the 24-month window (the cell's 6.9% rate) drop to zero by modelling assumption — pre-agreed terms eliminate the recoup arithmetic that every settlement-flow attention kind in this app traces back to.

PHASE 2 — \$1-5K VS AND % OF NET SIMPLIFICATION

Phase 2 • MVP

Convert vs and percentage-of-net deals to flat guarantees

137 vs • 102 % of net • 24 months

PROBLEM

A vs deal requires calculating gross revenue, applying the expense cap, computing a net base, running a percentage calculation, comparing it to the guarantee, and paying the higher number. Percentage-of-net deals involve the same arithmetic. Both structures exist to protect the artist when a show overperforms — when the percentage would produce more than the guarantee.

At The Crescent's gross levels this protection rarely engages. The expense structure consumes the available pool before the percentage produces anything meaningfully above the guarantee. The complexity is mostly cost, with a measurable price: vs deals in the \$1-5K bucket dispute at **12.4%** versus

3.8% for equivalent flat deals in the same bucket — more than three times the rate, on substantively similar economics.

FINDING

The cross-tab tells the story compactly:

DEAL TYPE · BUCKET	SHOWS	DISPUTE RATE	LOSS-MAKING RATE	ATTENTION ITEMS
flat · \$1–5K	104	3.8%	35.6%	4
vs · \$1–5K	137	12.4%	19.7%	16
vs · \$0–1K	14	21.4%	85.7%	3
% of net · \$0–1K	102	8.8%	36.3%	9

Switching just the 137 vs / \$1–5K deals to a flat would have eliminated 12 of the 17 disputes (12.4% → 3.8%) and saved the venue \$22,774 across the 24-month window — sourced directly from the projected-grid endpoint, where the actual sum of payouts (\$649,911) is compared to the Smart Switch flat sum (\$627,137). The dispute-reduction model assumes a 70% reduction (not 100%) — accepted flats can still surface disputes from soft-cost recoups or hospitality, just at the same baseline rate as historical flat deals in this bucket.

The disputes savings are not a free assumption. They follow from the structural change: a pre-agreed flat has nothing to argue about at midnight, and every settlement-flow attention rule in this app — disputed recoups, status-vs-notes mismatches, stale disputes — traces back to the recoup arithmetic that flats eliminate.

The percentage clause does fire — just unevenly. Across 180 settled vs deals in the window, the artist's percentage out-paid the guarantee in 156 (86.7%) of them; in the other 13.3% the artist was paid the floor exactly (avg guarantee-win = \$0). So the percentage clause is doing real work for artists most of the time — which means the case for switching to flat is not "the percentage produces nothing." It is "the dispute cost outweighs the artist-side benefit on average," and Smart Switch's cell-mean flat (\$23K savings on 137 deals) is what calibrates that trade-off honestly. Surfaced as a top-line tile on /insights via vsPercentageFiredStats .

SUGGESTION

Smart Switch generates a flat suggestion at deal entry — weeks before the show — using the cell-average artist payout for vs and percentage-of-net deals (rounded to the nearest \$50, gated on at least 3 historical comparable shows in the cell, demoted one tier for first-time artists at the venue). The agent confirms or declines before show night. If accepted, settlement becomes: artist gets \$X, wire Monday.

This is where the AI layer adds the most visible value. Instead of presenting a switch proposal as a form to review, the AI can present it conversationally: "Based on how 137 similar shows have settled here, a flat of \$1,500 puts the artist within \$200 of what the vs terms would have produced on 110 of those nights. Want to offer that now?" That is a Wednesday email, not a midnight calculation.

EXPECTED OUTCOME

Settlement for accepted conversions requires no calculation — the agreed flat is the payout. The 12.4% dispute rate on vs / \$1–5K deals drops to near zero for converted deals because there is nothing to dispute. The data Greenroom collects on converted shows is clean, consistent, and immediately useful for refining the next round of suggestions — without any data cleaning step.

COMBINED IMPACT · PHASES 1 + 2

IF SMART SWITCH HAD BEEN THE DEFAULT FOR THE LAST 24 MONTHS

Money saved · door + vs/\$1–5K combined	+\$73,244
Loss-making nights avoided	17 of 27 (63%)
Disputes avoided across the grid	19
Settlement-flow attention items avoided	18

Headline view in the running app: the four green tiles at the top of the "If Smart Switch + Improve Deal had been used" card on [/insights](#). Drill into any cell to see actual vs projected per show, with the formulas spelled out in the calculator-style explainers.

PHASE 3 — INSURANCE SERVICES

Phase 3 • Post-MVP

Settlement insurance as a platform service Optional • not required for Phases 1-2

WHAT IT DOES

Insurance is the safety net that makes committing to a pre-agreed settlement safe for both parties. Two independent products, priced per show at booking:

Product 1 — Instant Dispute Pay. When a settlement enters disputed status, Greenroom pays the resolution amount immediately — no back-and-forth. Recovered in the next billing cycle. Sized against The Crescent's actual data: **22 of 509 settlements (4.3%) were disputed in the last 24 months, with \$7,582 of disputed-recoup value across them — an average disputed amount of \$345 per disputed settlement.**

Product 2 — Suggestion Accuracy Pay. When the accepted flat suggestion differs from what the artist would have earned under the original deal by more than \$50, Greenroom pays the gap up to a cap. The cap has now been sized against the live SGP backtest ([/api/insights/guarantee-backtest](#)): across **321 scored deals**, the median absolute gap is **\$1,119**, with p75 = **\$1,939** and p90 = **\$3,344**. A \$400 cap would have fully covered only **25.2%** of historical gaps; for honest pricing the launch cap should be set at the median (\$1,100, ≈50% full coverage) or p75 (\$1,900, ≈75% full coverage), with the smaller cap pitched as partial coverage. The full coverage curve — \$100 →

6.9%, \$200 → 13.1%, \$400 → 25.2%, \$800 → 39.6%, \$1,500 → 64.2% — is now surfaced as a top-line tile on the SGP backtest section of [/insights](#).

WHAT IT COSTS EACH PARTY

Either party enrolls alone (one-sided) or both enroll together at a shared discount (two-sided). The recommended default for \$1–5K converted deals is Product 1 at the Basic tier. For door-deal conversions, Product 2 requires the Pro tier because the suggestion gap on door deals is materially larger.

Two-sided enrollment has a practical advantage beyond the price discount: when both parties have opted in, both have an incentive to accept the flat suggestion and resolve disagreements quickly. Greenroom's data estimates this reduces the expected claim rate by approximately 15%.

REVENUE FOR GREENROOM • PROPOSAL-SIDE MODELLING

Insurance creates a new revenue stream from existing activity. Greenroom has structural advantages over traditional insurers — no agents, no physical distribution, and a proprietary dataset that gives pricing precision unavailable to a general carrier. These compress the expense ratio toward zero, making a 50% gross margin realistic and conservative by industry standards.

PRODUCT 1 • BASIC TIER • 50% TARGET MARGIN • ALL FIGURES SOURCED FROM LIVE DISPUTE DATA

Expected cost / show · 4.3% rate × \$345 avg disputed	\$14.84
Suggested premium @ 50% margin	\$30 / show
Two-sided premium · each party	\$15 / show
Net margin per enrolled show	\$15.16
The Crescent · full adoption · 509 shows / 24 mo	~\$321 / mo
Platform · 300 venues · 30% adoption	~\$162K / yr
Platform · 300 venues · 50% adoption	~\$270K / yr

These figures are from Product 1 alone and assume no Product 2 enrollment. Pricing inputs are from the live Crescent dataset; the across-platform extrapolation assumes The Crescent's **4.3% headline dispute rate** is typical.

With no second-venue dataset available yet, the only honest sensitivity test is intra-venue: dispute rates across this dataset's cells span **3.8% (flat / \$1–5K)** to **12.4% (vs / \$1–5K)** and **21.4% (vs / \$0–1K)**. Bracketing platform extrapolation at **3.5% – 6.0%** (the headline $\pm$ a conservative band that excludes the most-disputed exotic cells, since simplified deals will under-represent them) gives an expected cost / show of **\$12.08 – \$20.70** versus the \$14.84 point estimate, and a platform-wide annual revenue band of roughly **\$130K – \$310K** versus the \$162K – \$270K point range — overlapping but wider on both sides. The direction of the economics is not sensitive to either band; at any adoption rate above 15% and any dispute rate inside this band, the product covers its expected claim costs from month one. This intra-venue band is a placeholder until a second venue's dataset lands; the right comparison is per-venue, not per-cell.

EXPECTED OUTCOME

Insurance is not required for the deal simplification to work — Phases 1 and 2 improve the settlement experience regardless. What insurance adds is confidence. An agent who might hesitate to commit to a flat deal weeks in advance is more likely to accept when they know that if the actual outcome differs meaningfully from the suggestion, the gap is covered. That confidence is what drives adoption of the simplified structure, and adoption of the simplified structure is what drives the data quality improvements that make the AI tools useful.

THE AI LAYER

Simplified deal structures are not just easier to settle — they are the prerequisite for AI tools that actually work. A system that needs to reconcile ambiguous expense categories, interpret handwritten cap agreements, and cross-reference multiple deal structures cannot be meaningfully automated. A system with clean, consistent flat deals and hybrid door terms can be — and the LLM enrichment pipeline already running on the Insights tab is a working proof of that: 237 of 558 settlements have machine-generated positive/negative summaries today, clustered into the dominant complaint themes per cell.

DEAL STRUCTURE SUGGESTION

At booking, the AI suggests the right deal type and flat amount based on the artist's history at the venue, the cell's expense patterns, and comparable shows. The agent sees a recommendation, not a blank form. This is what Smart Switch + Improve Deal already do today.

PRE-SHOW Q&A

"What would the artist earn if the show does \$4,000?" is a question that should have an instant answer on Wednesday, not a midnight calculation. With flat deals, the answer is always the same number.

AUTOMATED SETTLEMENT

A flat deal that was agreed at booking requires no settlement calculation. Greenroom can pre-fill the settlement form, flag it for one-click confirmation, and wire automatically — without anyone doing arithmetic.

DISPUTE PREVENTION

Most disputes here arise from ambiguous expense lines and recoup labels in vs / % of net structures. Flat deals eliminate the expense ambiguity entirely. The AI flags the small number of remaining edge cases before show night rather than after.

The measure of success is not better analysis of complex deals. It is fewer complex deals that need analysing. Every vs deal converted to a flat in Phase 2 is a settlement that requires no calculation, no review, and no dispute resolution. That is the compounding advantage of simplification — each clean deal makes the next one easier, and the data it generates makes the AI suggestions more reliable over time.

RECONCILIATION NOTES & NEXT STEPS

This revision (v2) reconciles the May 2026 draft against the live Greenroom data pipelines. Five categories of figure changed and four next-step items remain before this can be considered shippable to a venue partner.

WHAT CHANGED FROM V1

CLAIM IN V1	V1 FIGURE	LIVE DATA	SOURCE
Pale Lake actual venue net	-\$1,930	\$0 (door pays full pool	switch-savings · show_0207

CLAIM IN V1	V1 FIGURE	LIVE DATA	SOURCE
		above expenses)	
Pale Lake hybrid artist payout	\$9,474	\$10,020	switch-savings · counterfactual
Pale Lake hybrid venue net	\$5,983	+\$7,367	switch-savings · moneySavedToVenue
Door deals losing money	29 of 29 (100%)	27 of 29 (93.1%)	deal-analysis · door \$0-1K cell
vs deals in scope (Phase 2)	43	137 (\$1-5K) + 14 (\$0-1K)	deal-analysis cross-tab
vs / \$1-5K dispute rate	4.7%	12.4%	deal-analysis cross-tab
Settlement dispute rate (Phase 3 input)	2.724%	4.3% (22 / 509)	reports · disputedRate
Avg contested \$ per disputed settlement	\$384	\$345	reports · derived
Insurance Product 1 premium	\$21 one-sided / \$11 each two-sided	\$30 one-sided / \$15 each two-sided	repriced from live dispute data @ 50% margin

RECOMMENDED NEXT STEPS BEFORE PUBLISHING

- ✓ **Done — "% never fired" rollup on switch-savings.** Now exposed as a top-line tile (`vsPercentageFiredStats`) in the Switch Savings section of `/insights` . The honest answer reconciled here, scoped to the \$1-5K bucket where Smart Switch's flat fallback actually applies: the percentage clause out-paid the guarantee in **53 of 64 (82.8%)** settled vs / \$1-5K deals, with the remaining **17.2%** paid at the guarantee floor (avg guarantee-win \$0). The Phase-2 case for switching to flat is therefore the dispute-cost trade-off (12.4% disputes vs 3.8% on flat), not a claim that the percentage produces nothing.

2

✓ **Done — "suggestion gap $\leq$ \$T" rollup on the SGP backtest.** A `gapCoverage` tile now ships with the SGP backtest endpoint and is rendered above the per-deal table on `/insights`. Crescent distribution: median \$1,119, p75 \$1,939, p90 \$3,344. Phase-3 Product 2 cap sizing has been updated above with the live coverage curve.

3

✓ **Done (provisionally) — sensitivity band on platform extrapolation.** The Phase-3 revenue strip now reports a **\$130K – \$310K** intra-venue band around the \$162K–\$270K point figures, derived from per-cell dispute-rate spread (3.8% – 12.4%) and clipped to a 3.5% – 6.0% modelling range that excludes the most-disputed exotic cells (since simplified deals will under-represent them). Marked explicitly as a placeholder until a second venue's dataset lands — the right comparison is per-venue, not per-cell.

All three next-steps are now applied in the live app and reflected above. The structural argument — simplify the deal types Smart Switch already covers, then layer Improve-Deal caps and insurance on top — is fully supported by the live data as of this revision.