

Supporting Document — Drift Report

Generated: 2026-05-14 **Source of truth:** Live API at this commit **Audit target:** `exports/greenroom-supporting-document.{html,pdf}`

This document lists every cited figure in `greenroom-supporting-document` that no longer matches the live Greenroom API, with the proposed replacement text and the API call used to verify each number.

TL;DR

- **Stable (no change needed):** core Reports KPIs, the cross-tab cell counts and dispute rates (`137 vs/$1–5K deals` , `12.4% vs disputes` , `3.8% flat disputes`), and LLM enrichment coverage.
- **Drifted (substantive update needed):** every figure derived from the Smart Switch projected-grid and the SGP backtest distribution. The doc was generated against an older snapshot where the SGP backtest scanned ~321 deals; the current 12-month window scans 156. Six figures need to change.

Stable rows (verified, leave as-is)

Doc claim	Doc value	Live value	Notes
Past settlements	509	509	✓
Disputed settlements	22 (4.3%)	22 (4.32%)	✓
Disputed-recoup value (avg)	\$7,582 (\$345)	\$7,582 (\$344.6)	✓ rounding only
LLM enrichment coverage	237 / 558 (42%)	237 / 558 (42%)	✓
<code>vs · \$1–5K</code> deal count	137	137	✓
<code>vs · \$1–5K</code> dispute rate	12.4%	12.41%	✓
<code>flat · \$1–5K</code> dispute rate	3.8%	3.85%	✓

Endpoints used: `GET /api/reports` , `GET /api/deal-analysis`
(`revenue.crossTabBySizeAndType.cells`), `GET /api/insights` .

Drifted rows (need to be updated in the doc)

1. Smart Switch projected savings · vs/\$1-5K

- **Doc text:** "Switching the 137 vs/\$1-5K deals to flat would have saved \$23,024" (also rendered as \$22,774 in the supporting-numbers table).
- **Live value:** **\$9,186** savings to venue across **34** switch-eligible vs/\$1-5K deals.
- **Why it changed:** switch-projected-grid now applies stricter eligibility (must be settled + have all inputs available pre-show + Smart Switch must have a valid suggestion). Of the 137 vs/\$1-5K cells in the analytics grid, only 34 now satisfy that bar. The cell-count of 137 in the cross-tab is unaffected because the analytics grid is unconditional.
- **Endpoint:** GET /api/insights/switch-projected-grid , cells[dealType=vs, bucket=\$1-5K] .
- **Proposed replacement (table row):**

Smart Switch projected savings · vs/\$1-5K · \$9,186 across 34 switch-eligible deals

- **Proposed replacement (narrative line):**

Switching the 34 vs/\$1-5K deals Smart Switch could have handled would have saved \$9,186 (venue side). The wider 137-deal cell is the analytics base; the 34-deal slice is the actionable sub-set.

2. Door deals — largest savings cell

- **Doc text:** "Door deals are the single largest savings cell (~\$50K via hybrid)."
- **Live value:** Door cells combined now save **\$13,211** to the venue, concentrated in door · \$0-1K (6 deals, \$27,549 actual → \$14,338 projected = \$13,211 saved).
- **Why it changed:** door deals shifted out of larger size buckets as the data window slid; only the \$0-1K cell still has switch-applies coverage. The "\$50K" figure is no longer reproducible from the live grid.
- **Endpoint:** GET /api/insights/switch-projected-grid , sum of moneySavedToVenue over dealType="door" cells.
- **Proposed replacement:**

Door deals remain the highest per-deal savings shape (avg ~\$2,200 saved per converted door deal). The current 12-month window has 6 switch-eligible door deals concentrated in \$0-1K, for \$13,211 total venue-side savings.

3. vs deals · percentage clause never fired

- **Doc text:** "11/64 (17.2%) of settled vs / \$1-5K deals had the percentage clause never out-pay the guarantee."
- **Live value:** **17/133 (12.78%)** across the broader vs scan.
- **Why it changed:** `vsPercentageFiredStats` was re-scoped — the denominator widened from "\$1-5K only" to "all settled vs deals," so the rate dropped even though the absolute count of `never-fired` deals rose from 11 to 17. The complement statistic ("clause out-pays the guarantee") consequently changed too.
- **Endpoint:** `GET /api/insights/switch-savings , field vsPercentageFiredStats .`
- **Proposed replacement:**

17 of 133 (12.8%) settled vs deals had the percentage clause never out-pay the guarantee. The flat fallback is therefore the right primary recommendation for the cell.

- **Knock-on edit:** the line "The percentage clause out-pays the guarantee in 86.7% of vs deals (avg = \$0 when it doesn't)" should read **87.2%** (= 116/133). The "\$0 when it doesn't" half is still accurate (`avgGuaranteeWin = 0`).

4. SGP gap distribution (median · p75 · p90)

- **Doc text:** "median \$1,119 · p75 \$1,939 · p90 \$3,344"
- **Live value:** **median \$870 · p75 \$1,569 · p90 \$2,575.**
- **Why it changed:** the SGP backtest is now scoped to the rolling 12-month window (n=156). The doc's figures came from a wider window (n=321). The insurance pricing brief already supersedes these numbers explicitly — the supporting doc needs the same update.
- **Endpoint:** `GET /api/insights/guarantee-backtest , gapCoverage.{medianAbsDelta, p75AbsDelta, p90AbsDelta} .`
- **Proposed replacement:**

SGP gap distribution (median · p75 · p90) · \$870 · \$1,569 · \$2,575 (12-month window, n=156)

5. SGP gap fully covered at \$400

- **Doc text:** "25.2%"
- **Live value:** **29.5%** (46 of 156 deals fully covered at the \$400 threshold).
- **Why it changed:** smaller, more recent window → tighter typical gaps → higher coverage rate at any given cap.

- **Endpoint:** `GET /api/insights/guarantee-backtest , gapCoverage.buckets[threshold=400].rate .`
- **Proposed replacement:**

| SGP gap fully covered at \$400 · 29.5%

6. SGP backtest size used by Product 2 cap sizing

- **Doc text:** implicit in the description block: "Now also reports `gapCoverage` (median \$1,119 · p75 \$1,939 · p90 \$3,344) used for Phase-3 Product 2 cap sizing."
- **Live value:** the inline (median · p75 · p90) values are the ones flagged in row 4 above. Also: explicit n is not in the doc; the insurance brief now cites n=156 — for consistency, add this to the supporting doc too.
- **Proposed replacement:**

| Reports `gapCoverage` over the rolling 12-month window (n=156, median \$870 · p75 \$1,569 · p90 \$2,575); used for Phase-3 Product 2 cap sizing.

Document-level recommendations

1. **Add a regeneration timestamp** at the top of the doc (the insurance brief already does this). Without it, future readers can't tell at a glance whether numbers match the live API.
2. **Add a "regenerated against live API at commit X" line** so the reader knows which snapshot the doc is bound to.
3. **Optional consolidation:** consider moving the SGP-backtest figures into a single block that's regenerated from one API call. As written, the median/p75/p90 numbers appear in two separate places (the description block for `guaranteeBacktest` and the supporting-numbers table), making them easy to drift out of sync — as in fact happened.

Verification recipe

Anyone can reproduce this drift report with:

```
curl -s http://localhost:80/api/reports > /tmp/reports.json
curl -s http://localhost:80/api/deal-analysis > /tmp/da.json
curl -s http://localhost:80/api/insights > /tmp/ins.json
```

```
curl -s http://localhost:80/api/insights/guarantee-backtest > /tmp/gb.json  
curl -s http://localhost:80/api/insights/switch-savings > /tmp/ss.json  
curl -s http://localhost:80/api/insights/switch-projected-grid > /tmp/spg.json
```

then cross-reference each row above against the relevant JSON field.

Change count

- 7 stable rows (no edit needed).
- 6 drifted rows (substantive numeric updates).
- 1 supplementary line (avgGuaranteeWin complement: 86.7% → 87.2%).
- 0 schema changes — only the numbers move.